

Linear Regression

The Seatbelts data set contains the monthly totals of car drivers in Great Britain killed or seriously injured Jan 1969 to Dec 1984. Here is the general structure of Seatbelts data set.

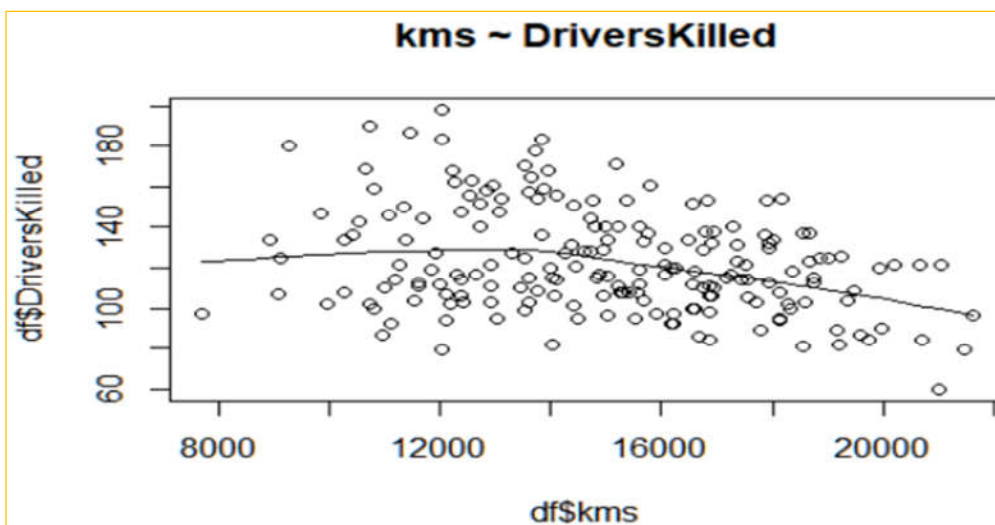
```
> data(Seatbelts)
> #View structure of data
> dat <- Seatbelts
> print(dat)
```

	DriversKilled	drivers	front	rear	kms	PetrolPrice	VanKilled	law
Jan 1969	107	1687	867	269	9059	0.1030	12	0
Feb 1969	97	1508	825	265	7685	0.1024	6	0
Mar 1969	102	1507	806	319	9963	0.1021	12	0
Apr 1969	87	1385	814	407	10955	0.1009	8	0
May 1969	119	1632	991	454	11823	0.1010	10	0
Jun 1969	106	1511	945	427	12391	0.1006	13	0
Jul 1969	110	1559	1004	522	13460	0.1038	11	0
Aug 1969	106	1630	1091	536	14055	0.1041	6	0
Sep 1969	107	1579	958	405	12106	0.1038	10	0
Oct 1969	134	1653	850	437	11372	0.1030	16	0
Nov 1969	147	2152	1109	434	9834	0.1027	13	0
Dec 1969	180	2148	1113	437	9267	0.1020	14	0
Jan 1970	125	1752	925	316	9130	0.1013	14	0

Building a simple regression model that we can use to predict DriversKilled by establishing a linear relationship with kms (Kilometers travelled). First we need to understand these variables graphically and visualize the following behavior:

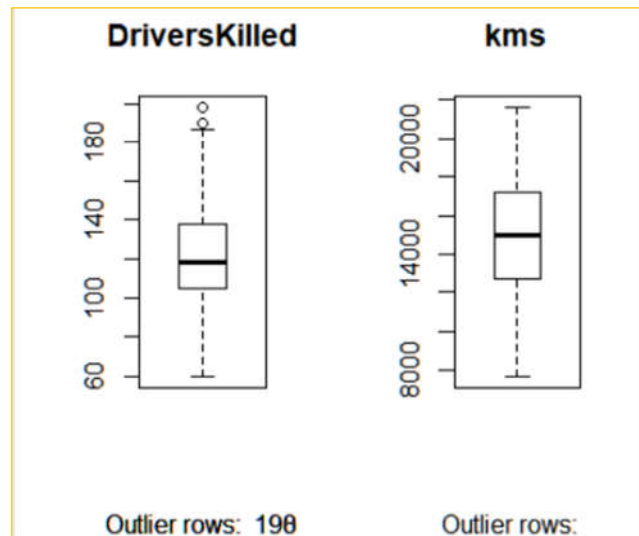
Scatter plot: Visualize the linear relationship between the predictor and response

```
> #Scatter Plot analysis
> scatter.smooth(x=df$kms, y=df$DriversKilled, main="kms ~ DriversKilled" )
> |
```



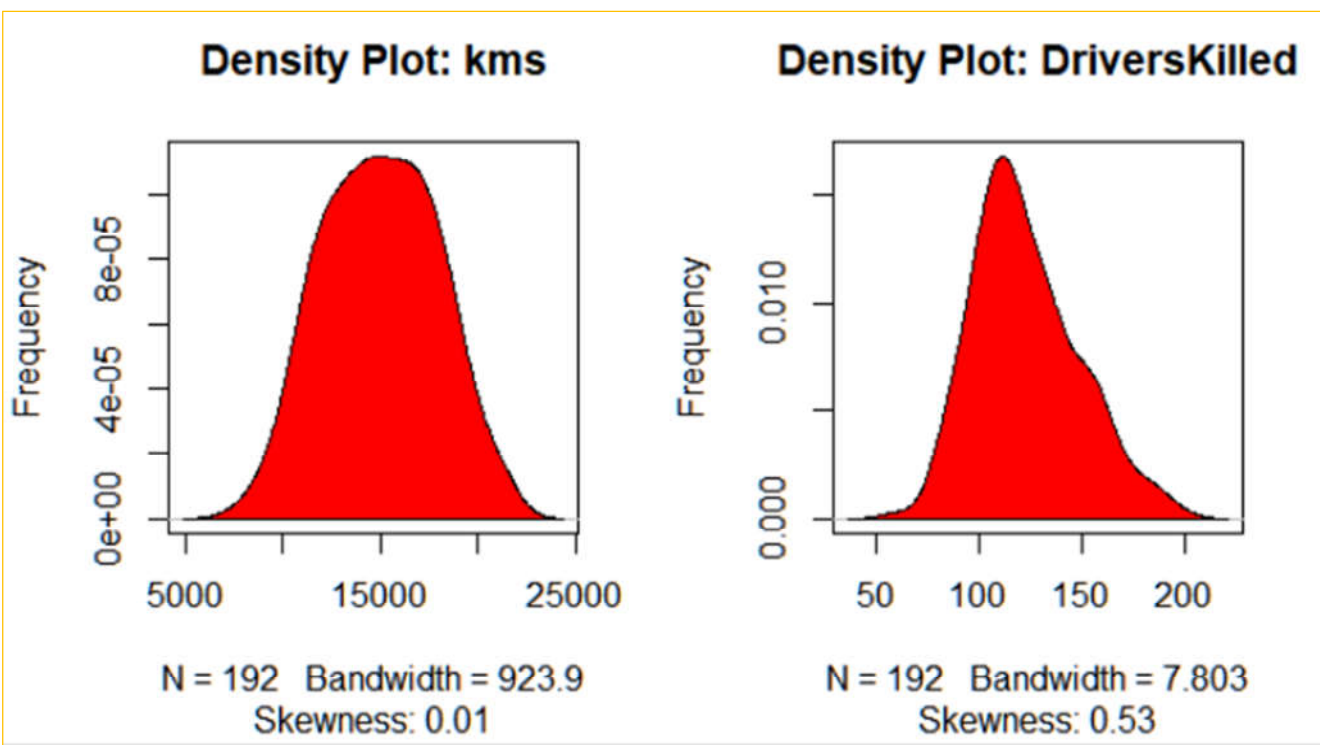
Box plot: To spot any outlier observations in the variable.

```
> #BoxPlot - Check for outliers
> par(mfrow=c(1, 2)) # divide graph area in 2 columns
> # box plot for 'DriversKilled'
> boxplot(df$DriversKilled, main="DriversKilled", sub=paste("Outlier rows: ",
+                                                         boxplot.stats(df$DriversKilled)$out))
> # box plot for 'kms'
> boxplot(df$kms, main="kms", sub=paste("Outlier rows: ", boxplot.stats(df$kms)$out))
> |
```



Density plot: To see the distribution of the predictor variable. Ideally, a close to normal distribution (a bell shaped curve), without being skewed to the left or right is preferred. Let us see how to make each one of them.

```
> #Density plot - Correlation
> # divide graph area in 2 columns
> par(mfrow=c(1, 2))
> # density plot for 'kms'
> plot(density(df$kms), main="Density Plot: kms",
+       ylab="Frequency", sub=paste("Skewness:", round(e1071::skewness(df$kms), 2)))
> polygon(density(df$kms), col="red")
> # density plot for 'DriversKilled'
> plot(density(df$DriversKilled), main="Density Plot: DriversKilled",
+       ylab="Frequency", sub=paste("Skewness:", round(e1071::skewness(df$DriversKilled), 2)))
> polygon(density(df$DriversKilled), col="red")
>
```



```
> ## calculate correlation between DriversKilled and kms
> cor(df$DriversKilled, df$kms)
[1] -0.3211016
```

Build Linear Model

```
> ## build linear regression model on full data
> linearMod <- lm(DriversKilled ~ kms, data=Seatbelts)
> print(linearMod)

Call:
lm(formula = DriversKilled ~ kms, data = Seatbelts)

Coefficients:
(Intercept)          kms
  164.391144    -0.002774
```

Now that we have built the linear model, we also have established the relationship between the predictor and response in the form of a mathematical formula for DriversKilled as a function for kms. For the above output, you can notice the ‘Coefficients’ part having two components: *Intercept*: 164.39, *kms*: -0.002774 these are also called the beta coefficients. In other words,

$$\text{DriversKilled} = \text{Intercept} + (\beta * \text{kms})$$

$$\text{DriversKilled} = 164.391 + (-0.00277 * \text{kms})$$

For example the number of kilometers travelled is 16,000 so according to the formula DriversKilled will be 120.

Linear Regression Diagnostics

```
> #Linear Regression Diagnostics
> summary(linearMod) # model summary

Call:
lm(formula = DriversKilled ~ kms, data = Seatbelts)

Residuals:
    Min       1Q   Median       3Q      Max
-52.028 -19.021  -1.974   16.719   66.964

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.644e+02  9.067e+00  18.130  < 2e-16 ***
kms          -2.774e-03  5.935e-04  -4.674  5.6e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 24.1 on 190 degrees of freedom
Multiple R-squared:  0.1031,    Adjusted R-squared:  0.09839
F-statistic: 21.84 on 1 and 190 DF, p-value: 5.596e-06
```

R-Squared and Adj R-Squared

```
> summary(linearMod) # model summary

Call:
lm(formula = DriversKilled ~ kms, data = Seatbelts)

Residuals:
    Min       1Q   Median       3Q      Max
-52.028 -19.021  -1.974   16.719   66.964

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.644e+02  9.067e+00  18.130 < 2e-16 ***
kms          -2.774e-03  5.935e-04  -4.674  5.6e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 24.1 on 190 degrees of freedom
Multiple R-squared:  0.1031,    Adjusted R-squared:  0.09839
F-statistic: 21.84 on 1 and 190 DF,  p-value: 5.596e-06
```

AIC and BIC

For model comparison, the model with the lowest AIC and BIC score is preferred.

```
> #Model comparision
> AIC(linearMod)
[1] 1770.816
> BIC(linearMod)
[1] 1780.589
> |
```

Predicting Linear Models

So far we have seen how to build a linear regression model using the whole dataset. If we build it that way, there is no way to tell how the model will perform with new data. So the preferred practice is to split your dataset into a 80:20 sample (training:test), then, build the model on the 80% sample and then use the model thus built to predict the dependent variable on test data.

Step 1: Create the training (development) and test (validation) data samples from original data.

```
> #Making predictions
> # Create Training and Test data -
> set.seed(100) # setting seed to reproduce results of random sampling
> trainingRowIndex <- sample(1:nrow(Seatbelts), 0.8*nrow(Seatbelts)) # row indices for training data
> trainingData <- Seatbelts[trainingRowIndex, ] # model training data
> testData <- Seatbelts[-trainingRowIndex, ] # test data
> |
```


Step 2: Develop the model on the training data and use it to predict the DriversKilled on test data

```
> #Making predictions
> # Step 1 Create Training and Test data -
> set.seed(100) # setting seed to reproduce results of random sampling
> trainingRowIndex <- sample(1:nrow(Seatbelts), 0.8*nrow(Seatbelts)) # row indices for training data
> trainingData <- as.data.frame.matrix(Seatbelts[trainingRowIndex, ]) # model training data
> testData <- as.data.frame.matrix(Seatbelts[-trainingRowIndex, ]) # test data
> #Step 2 Build the model on training data
> lmMod <- lm(DriversKilled ~ kms, data=trainingData) # build the model
> distPred <- predict(lmMod, testData) # predict distance
```

Step 3: Review diagnostic measures.

```
> #Step 3: Review diagnostic measures.
> summary (lmMod)

Call:
lm(formula = DriversKilled ~ kms, data = trainingData)

Residuals:
    Min       1Q   Median       3Q      Max
-52.110 -19.364  -2.137   16.959   66.881

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.654e+02  1.002e+01  16.504  < 2e-16 ***
kms          -2.853e-03  6.583e-04  -4.334  2.66e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 25.14 on 151 degrees of freedom
Multiple R-squared:  0.1106,    Adjusted R-squared:  0.1047
F-statistic: 18.78 on 1 and 151 DF,  p-value: 2.664e-05
```

From the model summary, the model p value and predictor's p value are less than the significance level, so we know we have a statistically significant model. Also, the R-Sq and Adj R-Sq are comparative to the original model built on full data.

Step 4: Calculate prediction accuracy and error rates

```
> #Step 4 Calculate prediction accuracy and error rates
> # make actuals_predicted dataframe.
> actuals_preds <- data.frame(cbind(actuals=testData$DriversKilled, predicted=distPred))
> correlation_accuracy <- cor(actuals_preds)
> head(actuals_preds)
  actuals predicteds
1     106    130.0743
2     103    128.5879
3     117    130.3881
4     157    126.6622
5     136    125.9860
6     140    129.1100
```

Now let's calculate the Min Max accuracy and MAPE:

```
> #Min Max accuracy and MAPE:
> min_max_accuracy <- mean(apply(actuals_preds, 1, min) / apply(actuals_preds, 1, max))
> print(min_max_accuracy) # min_max accuracy
[1] 0.8782031
> mape <- mean(abs((actuals_preds$predicted - actuals_preds$actuals))/actuals_preds$actuals)
> print(mape) # mean absolute percentage deviation
[1] 0.136092
> |
```

R CODE

```
#Install neccessary packages
install.packages("dplyr")
install.packages("e1071")
install.packages("DAAG")

#Import neccessary packages
library(dplyr)
library(e1071)
library(DAAG)

data(Seatbelts)
dat <-Seatbelts
#View structure of data
print(dat)

#Limit variables to two columns of interest
df <-data.frame(Seatbelts[,c("DriversKilled","kms")])

#Scatter Plot analysis
scatter.smooth(x=df$kms, y=df$DriversKilled, main="kms ~ DriversKilled" )

#BoxPlot - Check for outliers
par(mfrow=c(1, 2)) # divide graph area in 2 columns

#box plot for 'DriversKilled'
boxplot(df$DriversKilled, main="DriversKilled", sub=paste("Outlier rows: ",
boxplot.stats(df$DriversKilled)$out))

# box plot for 'kms'
boxplot(df$kms, main="kms", sub=paste("Outlier rows: ",
boxplot.stats(df$kms)$out))

#Density plot - Correlation
#divide graph area in 2 columns
par(mfrow=c(1, 2))
```

```

#density plot for 'kms'
plot(density(df$kms), main="Density Plot: kms",
      ylab="Frequency", sub=paste("Skewness:", round(e1071::skewness(df$kms), 2)))
polygon(density(df$kms), col="red")

#density plot for 'DriversKilled'
plot(density(df$DriversKilled), main="Density Plot: DriversKilled",
      ylab="Frequency", sub=paste("Skewness:",
round(e1071::skewness(df$DriversKilled), 2)))
polygon(density(df$DriversKilled), col="red")

#calculate correlation between DriversKilled and kms
cor(df$kms, df$DriversKilled)

#build linear regression model on full data
linearMod <- lm(DriversKilled ~ kms, data=Seatbelts)
print(linearMod)

#Linear Regression Diagnostics
summary(linearMod) # model summary

#The p Value: Checking for statistical significance
modelSummary <- summary(linearMod) # capture model summary as an object
modelCoeffs <- modelSummary$coefficients # model coefficients
print(modelCoeffs)
beta.estimate <- modelCoeffs["kms", "Estimate"] # get beta estimate for speed
std.error <- modelCoeffs["kms", "Std. Error"] # get std.error for speed
t_value <- beta.estimate/std.error # calc t statistic
p_value <- 2*pt(-abs(t_value), df=nrow(Seatbelts)-ncol(Seatbelts)) # calc p Value
f_statistic <- linearMod$fstatistic[1] # fstatistic
f <- summary(linearMod)$fstatistic # parameters for model p-value calc
model_p <- pf(f[1], f[2], f[3], lower=FALSE)

print(t_value)
print(p_value)
print(f_statistic)
print(model_p)

#AIC and BIC
AIC(linearMod)
BIC(linearMod)

```



```

#Making predictions
#Step 1 Create Training and Test data -
set.seed(100) # setting seed to reproduce results of random sampling
trainingRowIndex <- sample(1:nrow(Seatbelts), 0.8*nrow(Seatbelts)) # row indices
for training data
trainingData <- as.data.frame.matrix(Seatbelts[trainingRowIndex, ]) # model
training data
testData <- as.data.frame.matrix(Seatbelts[-trainingRowIndex, ]) # test data

#Step 2 Build the model on training data
lmMod <- lm(DriversKilled ~ kms, data=trainingData) # build the model
distPred <- predict(lmMod, testData) # predict distance

#Step 3: Review diagnostic measures.
summary (lmMod)

#Step 4 Calculate prediction accuracy and error rates
#make actuals_predicted data frame.
actuals_preds <- data.frame(cbind(actuals=testData$DriversKilled,
predicted=distPred))
correlation_accuracy <- cor(actuals_preds)
head(actuals_preds)

#Min Max accuracy and MAPE:
min_max_accuracy <- mean(apply(actuals_preds, 1, min) / apply(actuals_preds, 1,
max))
print(min_max_accuracy) # min_max accuracy
mape <- mean(abs((actuals_preds$predicted -
actuals_preds$actuals))/actuals_preds$actuals)
print(mape) # mean absolute percentage deviation

```